

[Get started](#)[Open in app](#)492K Followers · [About](#) [Follow](#)


You have **2** free member-only stories left this month. [Sign up for Medium and get an extra one](#)

Traffic Sign Recognition with TensorFlow 2.x

End to end example from raw image processing to model evaluation



Jun M. Jan 21 · 4 min read ★



Left: TensorFlow logo. Image credit: tensorflow.com Right: original and grayscale traffic sign from data below

Introduction

TensorFlow is a software library most famous for its flexibility and ease of use in neural networks. You can find a lot of examples online from image classification to object detection, but many of them are based on TensorFlow 1.x. It is a big change from TensorFlow 1.0 to 2.0 with a tighter Keras integration, where the focus is more on

higher level APIs. Many methods have been depreciated (or you may use `tf.compat.v1`). Model construction becomes a lot easier and default parameters in each model already work very well for general use. With all the benefits, it still provides flexibility should you need to change the parameters.

In this article, I will use TensorFlow 2.0 (more specifically, Keras in TensorFlow) to classify traffic signs. The dataset is available many places on the web, but I will use this one [hosted on Kaggle](#).

Data overview

The data package includes folders of `Train`, `Test` and a `test.csv`. There are a `meta.csv` and a `Meta` folder to show the standard image for each traffic sign. There is also a `signname.csv` for mapping a label to its description. `Train` folder contains 43 sub-folders whose names are the labels of the images in them. For example, all the images in folder `0` has a class label of `0` and so on... The images are of different sizes ranging from 20x20 to 70x70, and all have 3 channels: RGB.

So the first thing I have to do is to resize all the images to 32x32x3 and read them into a numpy array as training features. At the same time, I created another numpy array with labels of each image, which is from the fold name where the image loaded from.

```
In [2]: import cv2
import glob
import pickle
import numpy as np
import pandas as pd

# function to read and resize images, get labels and store
# them into np array
def get_image_label_resize(label, filelist, dim = (32, 32),
dataset = 'Train'):
    x = np.array([cv2.resize(cv2.imread(fname), dim, interp
olation = cv2.INTER_AREA) for fname in filelist])
    y = np.array([label] * len(filelist))

    #print('{} examples loaded for label {}'.format(x.shape
[0], label))
    return (x, y)

# data for label 0. I store them in parent level so that th
ey won't be uploaded to github
filelist = glob.glob('../Train/'+str(0)+'/*.png')
trainx, trainy = get_image_label_resize(0, glob.glob('../Tr
ain/'+str(0)+'/*.png'))
```

```
In [3]: # go through all others labels and store images into np arr
```

ay

load training and testing data.ipynb hosted with ♥ by GitHub

view raw

I need to do the same for testing images. However the labels for testing images are stored as `classId` in `test.csv` with paths of that image. So I use pandas to read the `csv` file, load the image from path and assign the corresponding `classId`.

From the training set, I randomly spitted 20% as validation set for use during the process of model training. The model accuracy of training and validation will give us information about underfitting or overfitting.

```
In [11]: # shuffle training data and split them into training and validation
indices = np.random.permutation(trainx.shape[0])
# 20% to val
split_idx = int(trainx.shape[0]*0.8)
train_idx, val_idx = indices[:split_idx], indices[split_idx:]
X_train, X_validation = trainx[train_idx,:], trainx[val_idx,:]
y_train, y_validation = trainy[train_idx], trainy[val_idx]
```

```
In [16]: # get overall stat of the whole dataset
n_train = X_train.shape[0]
n_validation = X_validation.shape[0]
n_test = X_test.shape[0]
image_shape = X_train[0].shape
n_classes = len(np.unique(y_train))
print("There are {} training examples ".format(n_train))
print("There are {} validation examples".format(n_validation))
print("There are {} testing examples".format(n_test))
print("Image data shape is {}".format(image_shape))
print("There are {} classes".format(n_classes))
```

```
There are 31367 training examples
There are 7842 validation examples
There are 12620 testing examples
```

create val.ipynb hosted with ♥ by GitHub

view raw

Next, I converted images to grayscale and normalized each pixels. Normalization makes model to converge more quickly.

```
In [15]: # convert the images to grayscale
X_train_gry = np.sum(X_train/3, axis=3, keepdims=True)
X_validation_gry = np.sum(X_validation/3, axis=3, keepdims=True)
X_test_gry = np.sum(X_test/3, axis=3, keepdims=True)
```

```
X_test_gry = np.sum(X_test/3, axis=3, keepdims=True)

# Normalize data
X_train_normalized_gry = (X_train_gry-128)/128
X_validation_normalized_gry = (X_validation_gry-128)/128
X_test_normalized_gry = (X_test_gry-128)/128
```

```
In [23]: import matplotlib.pyplot as plt
%matplotlib inline
plt.style.use('seaborn-colorblind')

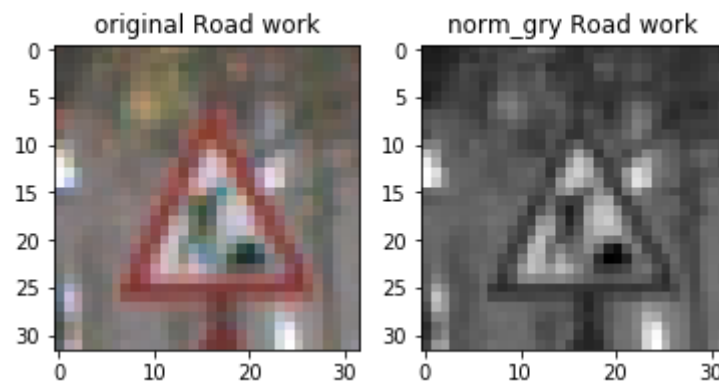
# descriptions for each label
sign = pd.read_csv('signnames.csv')

# pick an image, display the original and the normalized gray image
index = np.random.randint(0, n_train)
fig, ax = plt.subplots(1,2)
ax[0].set_title('original ' + sign.loc[sign['ClassId'] == y_train[index], 'SignName'].values[0])
ax[0].imshow(cv2.cvtColor(X_train[index], cv2.COLOR_BGR2RGB))
```

pre_processing.ipynb hosted with ♥ by GitHub

[view raw](#)

Here is a comparison between an RGB and grayscale image. The grayscale image still retains its features and can be recognized but with much smaller size.



A comparison of original and grayscale image

Model construction

I will use the famous LeNet published in 1998 by Yann LeCun et al. The input shape is 32x32x1. First convolution layer will have a depth of 6, a filter size of (5, 5), and a stride of (1, 1). Valid padding is used (i.e. no padding). Therefore the width (or height) of this layer is $32 - 5 + 1 = 28$, i.e. the shape is 28x28x6. The activation of this layer is relu.

```
In [78]: import tensorflow as tf
from tensorflow.keras import datasets, layers, models
```

```

model = models.Sequential()
# Conv 32x32x1 => 28x28x6.
model.add(layers.Conv2D(filters = 6, kernel_size = (5, 5),
                        strides=(1, 1), padding='valid',
                        activation='relu', data_format = 'c
hannels_last', input_shape = (32, 32, 1)))
# Maxpool 28x28x6 => 14x14x6
model.add(layers.MaxPooling2D((2, 2)))
# Conv 14x14x6 => 10x10x16
model.add(layers.Conv2D(16, (5, 5), activation='relu'))
# Maxpool 10x10x16 => 5x5x16
model.add(layers.MaxPooling2D((2, 2)))
# Flatten 5x5x16 => 400
model.add(layers.Flatten())
# Fully connected 400 => 120
model.add(layers.Dense(120, activation='relu'))
# Fully connected 120 => 84
model.add(layers.Dense(84, activation='relu'))
# Dropout
model.add(layers.Dropout(0.2))
# Fully connected, output layer 84 => 43
model.add(layers.Dense(43, activation='softmax'))

```

lenet.ipynb hosted with ♥ by GitHub

[view raw](#)

Following the first convolution layer is a max pooling layer. It effectively downsizes the data by only selecting the max value pixel for adjacent pixels. LeNet uses a (2, 2) kernel size. The default stride is the same as kernel, which means there is no overlap between the group of pixels the max is selected from. Now the shape of output becomes 14x14x6.

Next LeNet has a second convolution layer with depth of 16, filter size of (5, 5) and relu activation function, followed by a max pooling layer. The width (or height) of output is now $(14-5+1)/2 = 5$, i.e. the shape is 5x5x16.

The data is then flattened before the fully connected layers. The shape of output is $5 \times 5 \times 16 = 400$. This followed by 2 fully connected layers of size 120 and 84, with relu as activation function for both. A dropout layer is added to reduce overfitting. And finally a fully connected layer with size of 43 (the no. of classes). Softmax is used to return the probabilities of each class.

In [79]: `model.summary()`

Model: "sequential_7"

Layer (type)

Output Shape

Para

m #		
conv2d_12 (Conv2D)	(None, 28, 28, 6)	156
max_pooling2d_12 (MaxPooling)	(None, 14, 14, 6)	0
conv2d_13 (Conv2D)	(None, 10, 10, 16)	2416
max_pooling2d_13 (MaxPooling)	(None, 5, 5, 16)	0
flatten_6 (Flatten)	(None, 400)	0
dense_17 (Dense)	(None, 120)	4812

model summary.ipynb hosted with ♥ by GitHub

[view raw](#)

Model training and evaluation

Training is very straightforward with Keras. We only need to specify optimizer, loss function and validation metric. Within 10 epochs, the accuracy of both training and validation is above 0.97. With a dropout layer, there is no apparent overfitting. On the other hand, increasing training will only yield minimum improvement, so I stopped only after 10 epochs.

```
In [80]: # specify optimizer, loss function and metric
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

# training batch_size=128, epochs=10
conv = model.fit(X_train, y_train, batch_size=128, epochs=10,
                validation_data=(X_validation, y_validation))
```

```
Train on 31367 samples, validate on 7842 samples
Epoch 1/10
31367/31367 [=====] - 8s 255us/sample - loss: 2.2325 - accuracy: 0.4007 - val_loss: 0.8845 - val_accuracy: 0.7619
Epoch 2/10
31367/31367 [=====] - 8s 241us/sample - loss: 0.7074 - accuracy: 0.7961 - val_loss: 0.4006 - val_accuracy: 0.8986
Epoch 3/10
31367/31367 [=====] - 8s 242us/sample - loss: 0.3948 - accuracy: 0.8884 - val_loss: 0.2567 - val_accuracy: 0.9320
```

```


val_accuracy: 0.9242
Epoch 4/10
31367/31367 [=====] - 8s 243us/sample - loss: 0.2703 - accuracy: 0.9242 - val_loss: 0.1827 -

```

training.ipynb hosted with ♥ by GitHub

[view raw](#)

We can also plot the model performance on training and validation with each epoch. Indeed, the model appears to be quite generalized and not overfitting the training data.



Accuracy and loss of training and validation for each epoch

Finally, the model is used to predict the labels of test set. The accuracy is about 0.925.

```

12630/12630 [=====] - 1s 82us/sample - loss: 0.6084 - accuracy: 0.9250
Out[54]: [0.6084245350228081, 0.9250198]


```


```

In [97]: index = np.random.randint(0, n_test)
im = X_test[index]
fig, ax = plt.subplots()
ax.set_title(sign.loc[sign['ClassId'] == np.argmax(model.predict(np.array([im]))), 'SignName'].values[0])
ax.imshow(im.squeeze(), cmap = 'gray')

```

Out[97]: <matplotlib.image.AxesImage at 0x149dfdbd0>





test_pred.ipynb hosted with ❤ by GitHub

[view raw](#)

Conclusion

Previously I wrote an article on [building a neural network from scratch](#), which requires hardcore linear algebra. By using libraries like TensorFlow, the task becomes much easier and the model is more powerful.

You can get the full code from [here](#).

Sign up for The Daily Pick

By Towards Data Science

Hands-on real-world examples, research, tutorials, and cutting-edge techniques delivered Monday to Thursday. Make learning your daily ritual. [Take a look](#)

Your email

Get this newsletter

By signing up, you will create a Medium account if you don't already have one. Review our [Privacy Policy](#) for more information about our privacy practices.

Convolution Neural Net

Image Recognition

Python

TensorFlow

Keras